

Predicting nonlinear time series with reservoir computers

April 7, 2021

This document refers to exercise 9.9 in [1] and provides an example of how a reservoir computer can be implemented to predict a nonlinear time series. For a reference on the architecture and updating rules of reservoir computers, see pages 183 - 186 in [1]. The time series we are interested in predicting is given by the Lorenz equations

$$\begin{aligned}\frac{d}{dt}x_1 &= -\sigma x_1 + \sigma x_2, \\ \frac{d}{dt}x_2 &= -x_1 x_3 + r x_1 - x_2, \\ \frac{d}{dt}x_3 &= x_1 x_2 - b x_3\end{aligned}\tag{1}$$

which is a model for atmospheric convection. For the exercise, we are given the parameter values $\sigma = 10$, $r = 28$, and $b = 8/3$. Using these values, the system displays chaotic dynamics, which is characterised by (among other things) high sensitivity to initial conditions. The sensitivity can be quantified by considering the evolution of a small perturbation $\delta(t) \equiv |\delta(t)|$ from an initial state in phase space. One can show [2] that for large times t , the perturbation evolves as

$$\delta(t) \sim e^{\lambda_1 t} \delta(0).\tag{2}$$

The quantity λ_1 is the maximal Lyapunov exponent of the system. For chaotic systems, $\lambda_1 > 0$, and hence nearby trajectories diverge exponentially fast. The inverse λ_1^{-1} is called the Lyapunov time and gives an approximation of how long it takes before two nearby trajectories diverge significantly. For the Lorenz system with the provided parameter values, $\lambda_1 \approx 0.906$ [3].

Reservoir computers have been shown to be good at predicting nonlinear time series. In fig. 1, the dynamics of $x_2(t)$ is shown together with the prediction of a reservoir computer. As can be seen, the prediction remains accurate for more than twice the Lyapunov time. To obtain this result, eq. (1) is integrated numerically over $t = [0, 50]$ with time steps of 0.02, and initial conditions are set such that the trajectory starts on the chaotic attractor to avoid initial transients. The first 80% of the time series is then used for training, and the remaining 20% for testing. The input weights of the reservoir computer is initialised such that each entry is drawn from a uniform distribution ranging between $[-0.1, 0.1]$. The weights connecting the reservoir neurons is drawn from a uniform distribution ranging between $[-1, 1]$. The reservoir matrix is then re-scaled to make its maximal singular value is equal to unity (see page 185 of [1]). Finally, the weight matrix connecting the reservoir to the output is found through ridge regression with penalty parameter value 0.1, such that the difference between the output of the network and the target values are minimized. *tanh* is used as activation function for the reservoir neurons.

While the method described above yields decent results, it is by no means the only method, and other approaches may result in accurate time series predictions for longer times.

Training procedure

An example of the updating rules for a reservoir computer is

$$r_i(t+1) = g \left(\sum_j w_{ij} r_j(t) + \sum_{k=1}^N w_{ik}^{(in)} x_k(t) \right)\tag{3}$$

$$O_i(t+1) = \sum_{j=1}^M w_{ij}^{(out)} r_j(t+1)\tag{4}$$




Figure 1: Time series prediction of x_2 component of the Lorenz equations using a reservoir computer.

where $r_i(t)$ is the state of the i :th reservoir neuron at time t ($r_i(0) = 0$), $x_i(t)$ is the i :th input to the network at time t , w_{ij} are the weights connecting the reservoir neurons, $w_{ij}^{(in)}$ are the weights connecting the input to the reservoir, $w_{ij}^{(out)}$ are the weights connecting the reservoir to the output, $O_i(t)$ is the output at time t , and $g(x)$ is the activation function. To train a reservoir network, a time series is fed into the network and the states $r_i(0), r_i(1), \dots$ are recorded. Then, weights $w_{ij}^{(out)}$ are optimised (using e.g. least squares, gradient descent) such that the outputs O_i are as close to the desired outputs as possible. In our case, this would mean that the outputs give the coordinates (x_1, x_2, x_3) for the next time step in the time series. Once the network has been trained, the input $x_k(t)$ is replaced by $O_i(t)$ such that the network loops on itself. The output is then recorded for each time step.

References

- [1] B. Mehlig, *Artificial Neural Networks*, 2021, <https://arxiv.org/pdf/1901.05639.pdf>
- [2] K. Gustafsson, *Chaos and Lyapunov exponents*, 2017, <http://fy.chalmers.se/~f99krgu/dynsys/DynSysLecture10.pdf>
- [3] D. Viswanath, *Lyapunov exponents from random Fibonacci sequences to the Lorenz equations*, 1998, Cornell University